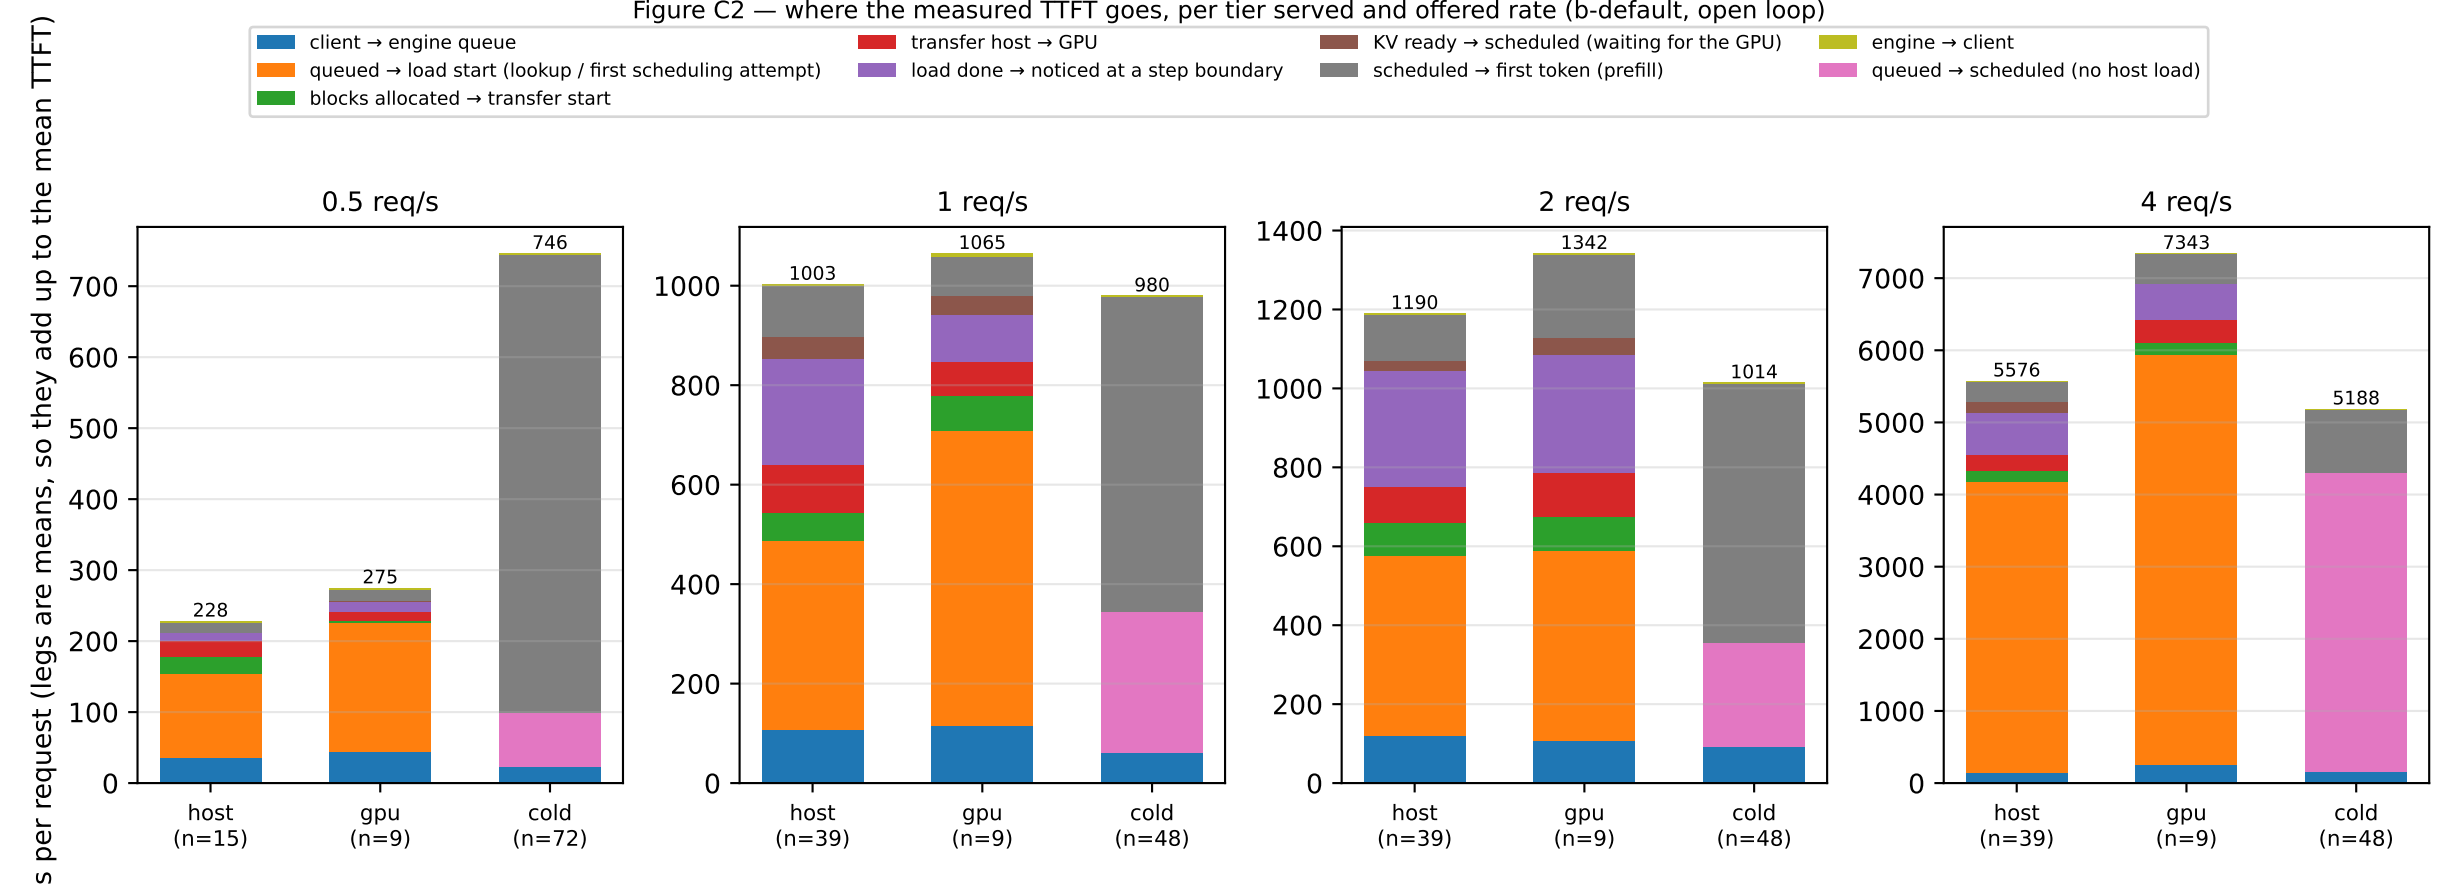
Figure C2 — where the measured TTFT goes, per tier served and offered rate (b-default, open loop)
client → engine queue
queued → load start (lookup / first scheduling attempt)
blocks allocated → transfer start
transfer host → GPU
load done → noticed at a step boundary
KV ready → scheduled (waiting for the GPU)
scheduled → first token (prefill)
engine → client
queued → scheduled (no host load)
s per request (legs are means, so they add up to the mean TTFT)
0.5 req/s
1 req/s
2 req/s
4 req/s
228
275
746
1003
1065
980
1190
1342
1014
5576
7343
5188
host (n=15)
gpu (n=9)
cold (n=72)
host (n=39)
gpu (n=9)
cold (n=48)
host (n=39)
gpu (n=9)
cold (n=48)
host (n=39)
gpu (n=9)
cold (n=48)